



Reasoning models solve the benchmark with fewer tokens each generation

OpenAI (GPT)



Anthropic (Opus)



● Mean tokens per problem ◆ Accuracy ■ Hard-but-doable problems (per-problem IQR)