

Accuracy and reasoning length across model generations (40 competition problems, right-censored at 32,768 tokens)





OpenAI (GPT)



Anthropic (Opus + Fable)



DeepSeek



All families censored at a common 32,768-token ceiling, so the rule is identical across columns -- but DeepSeek absorbs 23.6% of it against OpenAI's 1.4%, which biases its accuracy downward.
 DeepSeek = SiliconFlow/fp8 timeline (R1-0528 -> V3.2 -> V4 Pro April build).