



Output tokens: reasoning + answer (correct traces)

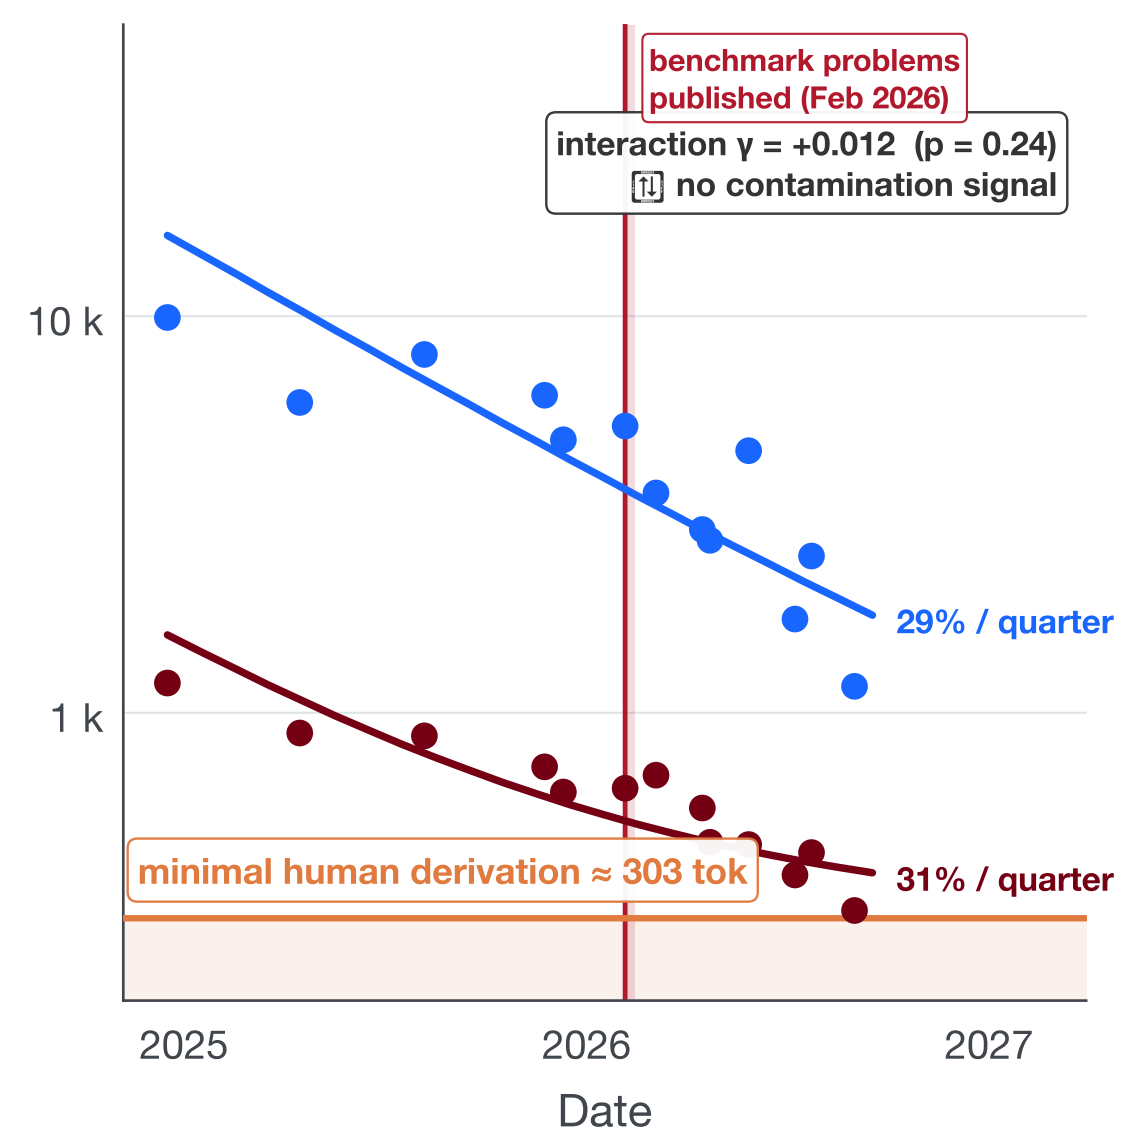
OpenAI (GPT) — pre-cutoff models only



GPT + Opus pooled — pre-cutoff models only



Same models, problems split by vintage
(full window — breaks the tie)



● Panels 1-2: pre-cutoff fit + models (cannot be contaminated) 95% CI (wild bootstrap) ● AIME/HMMT 2026 — new, only post-cutoff models could have seen (40 problems) Benchmark published (Feb 2026)

--- Pre-cutoff forecast - · - Full-sample fitted trend ● MATH-500 — old, in every model's training data (5 problems) Minimal human derivation