



DeepSeek: reasoning length and accuracy, with the V4 effort branch



max buys no accuracy on problems it finishes — it overruns the token budget ~40% more often. The branch measures the budget, not reasoning quality. April build offered high and max only (a requested low came back 0.5% from high). SiliconFlow fp8, censored at 32,768, 40 competition problems.