


# DeepSeek: reasoning length and accuracy, with the V4 effort branch



All models right-censored at 32,768 tokens: a trace needing more is scored wrong and clamped. OpenRouter/SiliconFlow fp8, 40 competition problems.  
April build offered high and max only — a requested low came back 0.5% from high, so it is not plotted.