



DeepSeek: reasoning length and accuracy, with the V4 effort branch

Capability keeps rising; accuracy under a fixed token budget does not



SOLID = accuracy under a 32,768-token budget (an overrun scores wrong). DASHED = accuracy among traces that FINISH, i.e. capability unconstrained by the budget. They diverge because V4 overruns the budget far more often: R1 and V3.2 were never cut at 32,768 by the provider (they ran to 66k and 83k), V4 was, on 41 of 320 trials. So V4 scoring below V3.2 on the solid line is a budget effect, not lower capability — on traces that finish V4 leads 98.8% to 92.7%. April build offered high and max only.